

(4)

The training error theorem for boosting

Here is pseudocode for the AdaBoost boosting algorithm presented in class:

Given: $(x_1, y_1), \dots, (x_N, y_N)$ where $x_i \in X$, $y_i \in \{-1, +1\}$

Initialize $D_1(i) = 1/N$.

For $t = 1, \dots, T$:

- Train weak learner using training data weighted according to distribution D_t .
- Get weak hypothesis $h_t : X \rightarrow \{-1, +1\}$.
- Measure "goodness" of h_t by its weighted error with respect to D_t :

$$\epsilon_t = \Pr_{i \sim D_t} [h_t(x_i) \neq y_i] = \sum_{i: h_t(x_i) \neq y_i} D_t(i).$$

- Let $\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$.

- Update:

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } y_i = h_t(x_i) \\ e^{\alpha_t} & \text{if } y_i \neq h_t(x_i) \end{cases} \quad (1)$$

where Z_t is a normalization factor (chosen so that D_{t+1} will be a distribution).

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right).$$

Although the notation is different, this algorithm is the same as in Fig. 18.10 of R&N.

In this note, we prove the training error theorem, which states that the training error of H is at most

$$\exp \left(-2 \sum_{t=1}^T \gamma_t^2 \right)$$

where $\epsilon_t = \frac{1}{2} - \gamma_t$.

We prove this in three steps.

Step 1: The first step is to show that

$$D_{T+1}(i) = \frac{1}{N} \cdot \frac{\exp(-y_i f(x_i))}{\prod_t Z_t}$$

where

$$f(x) = \sum_t \alpha_t h_t(x).$$

Proof: Note that Eq. (1) can be rewritten as

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$



since y_i and $h_t(x_i)$ are both in $\{-1, +1\}$. Unwrapping this recurrence, we get that

$$\begin{aligned} D_{T+1}(i) &= D_1(i) \cdot \frac{\exp(-\alpha_1 y_i h_1(x_i))}{Z_1} \dots \frac{\exp(-\alpha_T y_i h_T(x_i))}{Z_T} \\ &= \frac{1}{N} \cdot \frac{\exp(-y_i \sum_t \alpha_t h_t(x_i))}{\prod_t Z_t} \\ &= \frac{1}{N} \cdot \frac{\exp(-y_i f(x_i))}{\prod_t Z_t}. \end{aligned}$$

Step 2: Next, we show that the training error of the final classifier H is at most

$$\prod_{t=1}^T Z_t.$$

Proof:

$$\begin{aligned} \text{training error}(H) &= \frac{1}{N} \sum_i \begin{cases} 1 & \text{if } y_i \neq H(x_i) \\ 0 & \text{else} \end{cases} && \text{by definition of the training error} \\ &= \frac{1}{N} \sum_i \begin{cases} 1 & \text{if } y_i f(x_i) \leq 0 \\ 0 & \text{else} \end{cases} && \text{since } H(x) = \text{sign}(f(x)) \text{ and } y_i \in \{-1, +1\} \\ &\leq \frac{1}{N} \sum_i \exp(-y_i f(x_i)) && \text{since } e^{-z} \geq 1 \text{ if } z \leq 0 \\ &= \sum_i D_{T+1}(i) \prod_t Z_t && \text{by Step 1 above} \\ &= \prod_t Z_t && \text{since } D_{T+1} \text{ is a distribution} \end{aligned}$$

Step 3: The last step is to compute Z_t .

We can compute this normalization constant as follows:

$$\begin{aligned} Z_t &= \sum_i D_t(i) \times \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = y_i \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq y_i \end{cases} \\ &= \sum_{i: h_t(x_i)=y_i} D_t(i) e^{-\alpha_t} + \sum_{i: h_t(x_i) \neq y_i} D_t(i) e^{\alpha_t} \\ &= e^{-\alpha_t} \sum_{i: h_t(x_i)=y_i} D_t(i) + e^{\alpha_t} \sum_{i: h_t(x_i) \neq y_i} D_t(i) \\ &= e^{-\alpha_t} (1 - \epsilon_t) + e^{\alpha_t} \epsilon_t && \text{by definition of } \epsilon_t \\ &= 2\sqrt{\epsilon_t(1 - \epsilon_t)} && \text{by our choice of } \alpha_t \text{ (which was chosen to minimize this expression)} \\ &= \sqrt{1 - 4\gamma_t^2} && \text{plugging in } \epsilon_t = \frac{1}{2} - \gamma_t \\ &\leq e^{-2\gamma_t^2} && \text{using } 1 + x \leq e^x \text{ for all real } x \end{aligned}$$

Combining with Step 2 gives the claimed upper bound on the training error of H .

